

MIRA SOLENNE

Houston, TX | 713-555-0142 | mira.solenne@example.com | github.com/mira-solenne-labs

Professional Summary

AI engineer with expertise in RAG, production AI systems, and scalable data pipelines. Specialized in building retrieval-augmented knowledge APIs and deploying AI models in enterprise environments.

Education

Gulf Prairie Institute of Technology <i>M.S. in Applied Data Science, GPA 3.82/4.00</i>	Houston, TX December 2023
Redwood Plains University <i>B.S. in Computer Science</i>	San Marcos, TX May 2021

Experience

Machine Learning Engineer <i>Northstar Lantern Analytics</i>	August 2023 – Present Houston, TX
--	--------------------------------------

- Built Python and scikit-learn risk models over SQL data pipelines, raising validated recall by 14% while holding the false-positive rate constant.
- Delivered a retrieval-augmented knowledge API with FastAPI, embeddings, and vector search, reducing median analyst lookup time by 36% in a controlled pilot.
- Containerized models on AWS with MLflow, pytest, and drift alerts, cutting failed monthly releases from five to two in six months.

Data Science Intern <i>Copper Finch Systems</i>	June 2022 – August 2022 Austin, TX
---	---------------------------------------

- Prepared pandas and NumPy sensor features and trained a PyTorch anomaly model that improved F1 by 9% over the baseline.
- Created SQL quality checks and Git-based test fixtures, reducing recurring validation defects by 28% across four reporting runs.

Projects

CarePath Evidence Navigator <i>Python, FastAPI, PostgreSQL, Embeddings, Vector Search, Docker</i>	2025 Healthcare
---	--------------------

- Retrieved a reviewer-approved supporting passage in the top three results for 87% of 120 synthetic evaluation questions.

ShelfSense Demand Forecaster <i>Python, SQL, scikit-learn, pandas, MLflow</i>	2024 Retail and Supply Chain
---	---------------------------------

- Created a weekly forecasting pipeline that reduced holdout MAPE from 24.8% to 18.6% on synthetic store and product histories.

ModelWatch Release Guard <i>Python, Bash, AWS, Docker, MLflow, pytest, CI/CD</i>	2025 MLOps
--	---------------

- Detected all eight seeded regressions and reduced repeatable manual release checks from 42 to 15 minutes.

Skills

- **Languages:** Python, SQL, Bash
- **ML, Data & GenAI:** scikit-learn, PyTorch, pandas, NumPy, NLP, Transformers, Time Series Forecasting, Computer Vision, Model Evaluation, Retrieval-Augmented Generation, Embeddings, Vector Search, Prompt Engineering
- **Cloud, MLOps & Systems:** AWS, MLOps, Docker, MLflow, Model Monitoring, CI/CD, FastAPI, REST APIs, Git, pytest, PostgreSQL